

# Database Implementation and Indexing

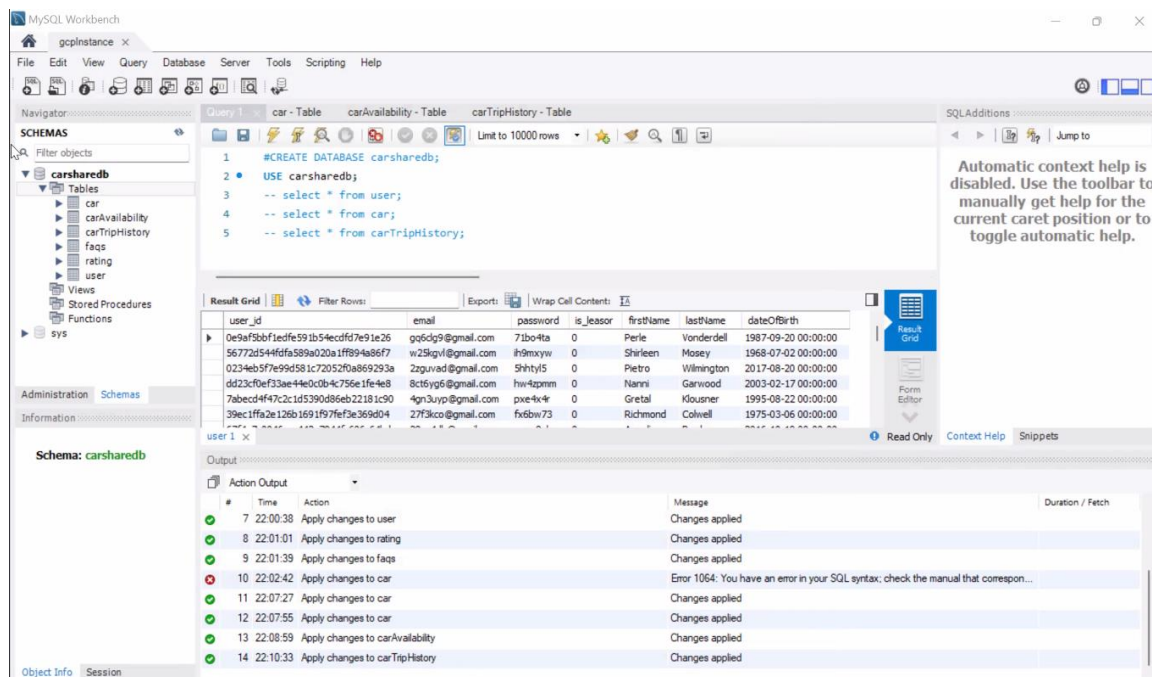
## Database Implementation

Our group used MySQL Workbench to ease the workflow of uploading our data and to make uploading quick and efficient.

The main tables which contain critical application information that we implemented are:

1. Car (Stores information such as make, model, year, VIN number)
2. CarAvailability (Stores information on when the car is available to rent)
3. CarTripHistory (Stores information including the number of trips the car has taken)
4. FAQs (Stores questions and answers for the user to view)
5. Ratings (Stores ratings for each car)

We followed MySQL workbench in order to upload our data. The instructions we used in order to upload this data are in the form of images below:



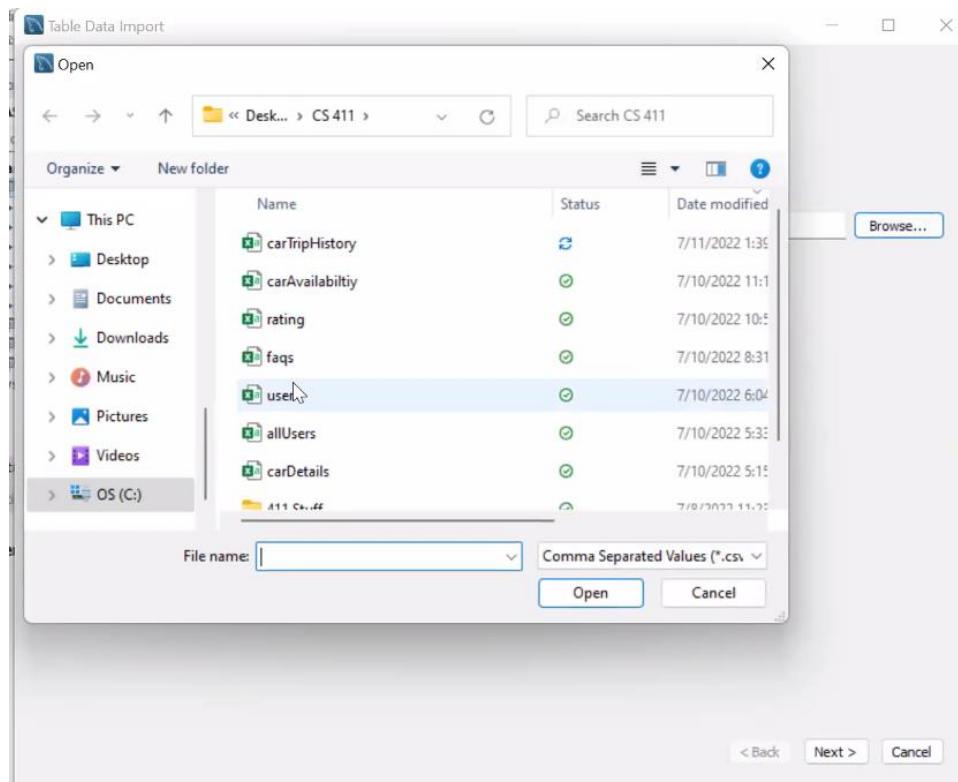
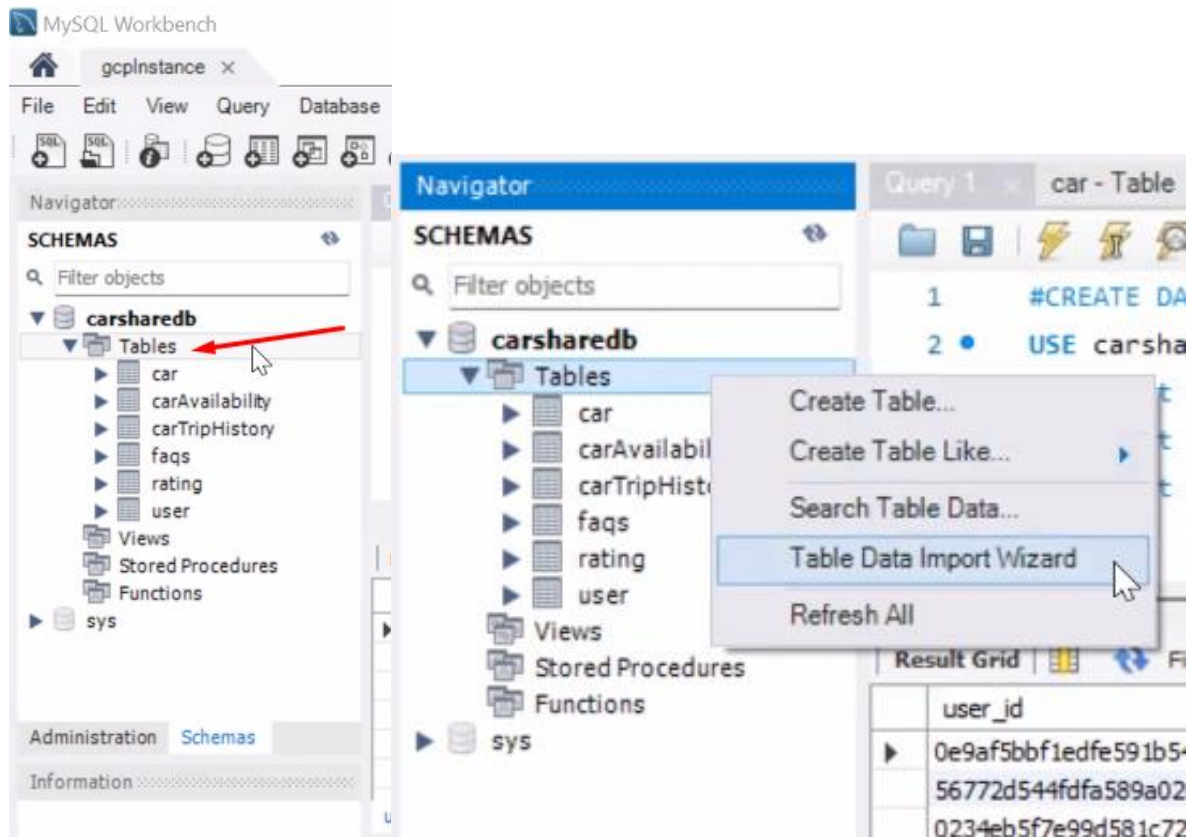


Table Data Import

Select File to Import

Table Data Import allows you to easily import CSV, JSON datafiles.  
You can also create destination table on the fly.

File Path: C:\Users\paary\OneDrive - University of Illinois - Urbana\Desktop\CS 411\user.csv

Browse...

Table Data Import

Select Destination

Select destination table and additional options.

☐ Use existing table: carsharedb.car

☒ Create new table: carsharedb , user

☒ Drop table if exists

< Back

Next >

Cancel

Table Data Import

Configure Import Settings

Detected file format: csv

Encoding: utf-8

Columns:

☒ Source Column

Field Type

user\_idtext

emailtext

passwordtext

is\_leaseorint

firstNametext

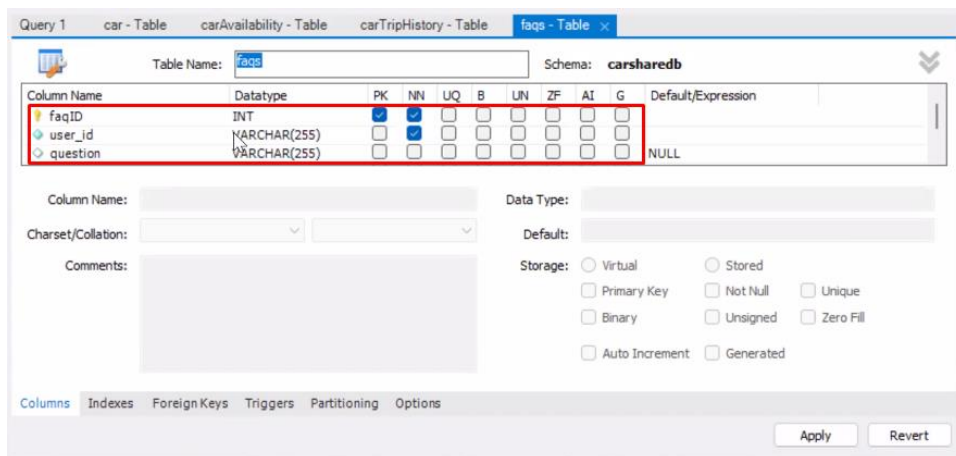
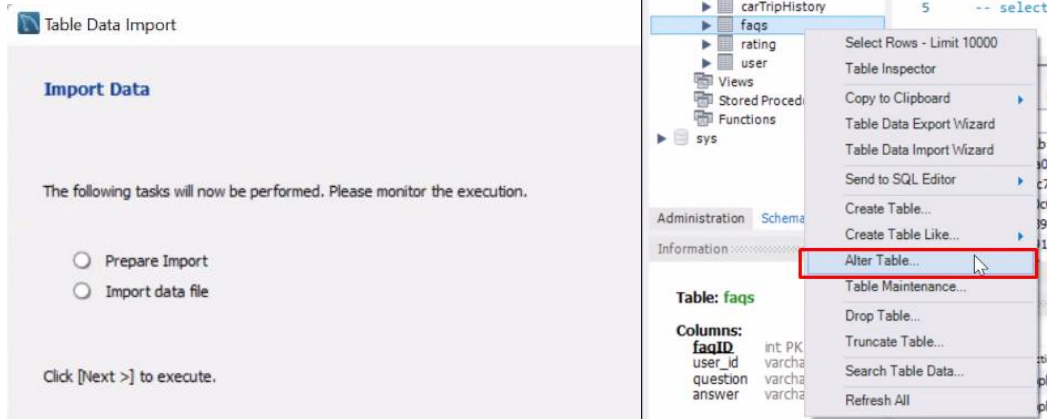
lastNametext

user_id	email	password	is_leaseor	firstName	lastName	dateOfBirth
0e9af5bbf...	gq6clg9@g...	71bo4ta	0	Perle	Vonderdell	1987-09-20
56772d54...	w25kgvl@g...	ih9mxyw	0	Shirleen	Mosey	1968-07-02
0234eb5f...	Zzquvad@g...	5hhtyl5	0	Pietro	Wilmington	2017-08-20
dd23cf0ef...	8ctgyg6@g...	hw4zpm	0	Nanni	Garwood	2003-02-17
7abec44f...	4gn3uyp@...	pxe4x4r	0	Gretal	Klousner	1995-08-22

< Back

Next >

Cancel



The **SQL DDL** commands we would use otherwise to create our tables would be as follows:

```
CREATE TABLE Car(car_id VARCHAR(8), car_lat DECIMAL(3,10), car_long DECIMAL(3,10), make VARCHAR(255), model VARCHAR(255), year INTEGER(4), vin VARCHAR(255));
```

```
CREATE TABLE CarTripHistory(journey_id VARCHAR(255), lesser_id VARCHAR(255), car_id VARCHAR(255), pickup_datetime DATETIME, dropoff_datetime DATETIME, price INTEGER(8), rating INTEGER(8));
```

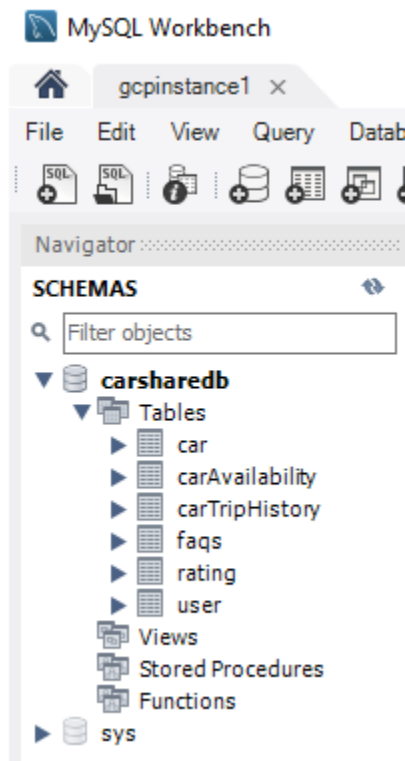

```
CREATE TABLE Caravailability(availability_id VARCHAR(255), car_id VARCHAR(255), available_from DATETIME, available_till DATETIME);
```

```
CREATE TABLE FAQs(faq_ID VARCHAR(255), user_id VARCHAR(255), question VARCHAR(255), answer VARCHAR(255));
```

```
CREATE TABLE Rating(rating INTEGER(8), message VARCHAR(255));
```

**Screenshot of connection to MySQL workbench/Google GCP:** We created a google GCP instance and then we connected to it through MySQL workbench. We then created a database called “carsharedb” to store all our data.

## MySQL Connections



Car Table Count

10 •	SELECT COUNT(*) from car;
11	#SELECT COUNT(*) from carAvailability;
12	#SELECT COUNT(*) from carTripHistory;
<	
Result Grid    Filter Rows: <input type="text"/>   Export:	
	COUNT(*)
▶	11706

Car Availability Count

11 •	SELECT COUNT(*) from carAvailability;
12	#SELECT COUNT(*) from carTripHistory;
13	
<	
Result Grid    Filter Rows: <input type="text"/>   Export:	
	COUNT(*)
▶	11706

CarTripHistory Table Count

12 •	SELECT COUNT(*) from carTripHistory;
13	
<	
Result Grid    Filter Rows: <input type="text"/>   Export:  Wrap Cell Content:	
	COUNT(*)
▶	4076

Rating Table Count

13 •	SELECT COUNT(*) from rating;
14	#####
<	
Result Grid    Filter Rows: <input type="text"/>   Export:	
	COUNT(*)
▶	4076

## Two Advanced SQL Queries:

**Query 1:** (Car location details with > 4.0 rating and > 50 trips)

```

7   # Finding car location details of those car which has more than 4.0 average user ratings
8   # and had more than 50 trips
9   • USE carsharedb;
10  • SELECT
11      car.car_id,
12      car_lat,
13      car_lon,
14      AVG(rating) AS avg_rating,
15      COUNT(*) AS total_trips
16  FROM car
17  JOIN carTripHistory ON car.car_id = carTripHistory.car_id
18  GROUP BY car.car_id
19  HAVING avg_rating >=4 AND total_trips >=50
20  ORDER BY total_trips DESC, avg_rating DESC
21  LIMIT 15;

```

Top 15 Results:

car_id	car_lat	car_lon	avg_rating	total_trips
406921adcca37705ef527b4246c10d2c	-12.1171111	-77.0389639	4.9974	459
0accdd3aa5a322f4129fa20b53278c69	-12.0918374	-77.0421799	4.9903	356
e84fda9c5df33f03c89b6923c362193f	-12.098176	-76.971417	4.9840	280
d04f1a596fe4a582f48e6eb2f9e8cceb	-12.024053	-77.112036	4.3365	241
35a17a45aa341b999787fa5e50608f65	-12.105472	-77.018854	4.8909	188
ff5c924e0b630fd7c019a423405c86f6	-12.098176	-76.971417	4.7338	174
f622a9397b85b5c9c0ef938c5d592020	-12.105472	-77.018854	4.9856	164
294e869d187357cad25b8af65aa860e8	-12.1053457	-76.9757211	5.0000	154
baacf396f773709519bbde35a585d91b	-12.1053457	-76.9757211	4.0000	74
c31d580a863a6601a97b1b12898445d4	-12.105472	-77.018854	5.0000	70
fb81277200c0aea311ad85f90d733b00	-12.024053	-77.112036	5.0000	67
922d998685f68f17dd0f3c9f36dd51ac	-12.105472	-77.018854	4.8868	64
62c2e39787aa19cd2513901cf3b49643	-12.024053	-77.112036	4.6724	64
b12f4f09c783e29fe0d0ea624530db56	-12.105472	-77.018854	4.9649	61
2a7979f6826bbbf4a061456c47623b7c	-12.105472	-77.018854	5.0000	58

Justification: This query will be a critical part of our website. What the query does is it returns important car information (such as its ID and location) for cars that meet certain requirements. These requirements are cars that have a rating that is greater than 4.0 and that have more than 50 trips. Our criteria for the thresholds are tailored so that the user has an easier time searching for cars that are reliable and that have a high rating. We decided that cars that have 50 trips or more while still holding an average rating that is greater than 4.0 would be the cars that most users would want to rent, so this query will assist in the user's search.

**Query 2:** (Show car availability of cars without any trip history)

```

23  # Finding out available cars in user defined timeframe
24  # which did not have any trip history (unused/new cars)
25  • USE carsharedb;
26  • SELECT
27      car.car_id, make, model, year, availability_from, availability_till
28  FROM carAvailability
29  JOIN car ON car.car_id = carAvailability.car_id
30  WHERE availability_from >= CURDATE() AND availability_till <=ADDDATE(CURDATE(), 5)
31      AND car.car_id NOT IN (SELECT DISTINCT car_id FROM carTripHistory)
32  ORDER BY availability_from ASC
33  LIMIT 15;

```



Top 15 Results:

car_id	make	model	year	availability_from	availability_till
a70d782602d74fec18457a2f4dba07	Audi	90	1988	2022-07-12 01:30:00	2022-07-12 17:00:00
4851ee4c867c9b49b909c1adb76c538e	Dodge	Ram 1500 Club	1997	2022-07-12 02:30:00	2022-07-12 19:00:00
54c88ff5110a79082d341066aa44761f	Mazda	626	1991	2022-07-12 02:30:00	2022-07-12 19:00:00
b3ba9cda1b124b4fdc5d639f2416c61c	Buick	Terraza	2007	2022-07-12 03:30:00	2022-07-12 21:30:00
b0f0845b5094fda2fcdac195eef9327	Infiniti	QX56	2012	2022-07-12 03:30:00	2022-07-12 21:00:00
bc5c1ae63a528f0371154594b34877d5	Honda	Odyssey	1999	2022-07-12 03:30:00	2022-07-12 21:00:00
d2cf43d9f47c6512b84e1ea7b87503d7	Mercury	Villager	1994	2022-07-12 14:30:00	2022-07-13 19:30:00
406921adcca37705ef527b4246c1a5c6	Dodge	Daytona	1992	2022-07-12 18:30:00	2022-07-13 06:00:00
2b9be9ff113177942788e7d6e85526b9	Dodge	Dakota Club	2002	2022-07-12 20:30:00	2022-07-14 07:00:00
1cbe459130a435037ebacaf7fb21861e	Volvo	S70	2000	2022-07-12 20:30:00	2022-07-14 07:30:00
12c472dd8efc6ed22ae866925ddd065	Ford	Expedition	2001	2022-07-13 03:30:00	2022-07-14 21:30:00
b1a3f957deba8f50052700bfb28a496a	Mercury	Monterey	2004	2022-07-13 03:30:00	2022-07-14 21:00:00
478419e53317b50cbf9af6b2e35ec463	Honda	Insight	2001	2022-07-13 06:30:00	2022-07-15 03:30:00
0ed064de66b8552b422f677d2c84462	Hummer	H2	2003	2022-07-13 08:00:00	2022-07-14 17:30:00
af07f9929390d66edf63ab7ee4406a94	Saturn	Sky	2009	2022-07-13 08:30:00	2022-07-15 07:30:00

Justification: This query is similar to the previous query, but it returns cars without any trip history associated with it in a certain timeframe. This will be another important piece on our website that will be used extensively if the user wants to rent a car that has not been used before in terms of renting.

## Indexing Analysis:

### Query 2:

- Default Index: car\_id, availability\_id
  - Time to execute: 0.02 sec
  - Query Cost:

```

-> Sort: total_trips DESC, avg_rating DESC (actual time=14.560..14.563 rows=18 loops=1)
-> Filter: ((avg_rating >= 4) and (total_trips >= 50)) (actual time=14.453..14.526 rows=18 loops=1)
-> Table scan on <temporary> (actual time=0.002..0.025 rows=135 loops=1)
-> Aggregate using temporary table (actual time=14.442..14.476 rows=135 loops=1)
-> Nested loop inner join (cost=1863.10 rows=4123) (actual time=0.080..10.391 rows=4076 loops=1)
-> Table scan on carTripHistory (cost=420.05 rows=4123) (actual time=0.055..1.688 rows=4076 loops=1)
-> Single-row index lookup on car using PRIMARY (car_id=carTripHistory.car_id) (cost=0.25 rows=1) (a
ctual time=0.002..0.002 rows=1 loops=4076)
1 row in set (0.02 sec)
mysql>

```

	Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null
►	carAvailability	0	PRIMARY	1	availability_id	A	11485			
◄	carAvailability	1	car_id_idx	1	car_id	A	126			

Query cost: 3856.09

query\_block #1

- Index Design #1: car\_id, availability\_id, availability\_from
  - Time to execute: 0.01 sec



	Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed
▶	carAvailability	0	PRIMARY	1	availability_id	A	11485	<b>NULL</b>	<b>NULL</b>
	carAvailability	1	car_id_idx	1	car_id	A	126	<b>NULL</b>	<b>NULL</b>
	carAvailability	1	idx_avail_from	1	availability_from	A	8393	<b>NULL</b>	<b>NULL</b>

query\_block #1

- | Table           | Non_unique | Key_name | Seq_in_index | Column_name       | Collation | Cardinality | Sub_part | Packed | Null |
|-----------------|------------|----------|--------------|-------------------|-----------|-------------|----------|--------|------|
| carAvailability | 0          | PRIMARY  | 1            | availability_id   | A         | 11485       |          |        |      |
| carAvailability | 1          |          | 1            | car_id            | A         | 126         |          |        |      |
| carAvailability | 1          |          | 1            | availability_till | A         | 8394        |          |        |      |

query\_block #1

- | Result Grid |                 | Filter Rows: | Export:        | Wrap Cell Contents: |                   |           |             |             |             |
|-------------|-----------------|--------------|----------------|---------------------|-------------------|-----------|-------------|-------------|-------------|
|             | Table           | Non_unique   | Key_name       | Seq_in_index        | Column_name       | Collation | Cardinality | Sub_part    | Packed      |
| ▶           | carAvailability | 0            | PRIMARY        | 1                   | availability_id   | A         | 11485       | <b>NULL</b> | <b>NULL</b> |
|             | carAvailability | 1            | car_id_idx     | 1                   | car_id            | A         | 126         | <b>NULL</b> | <b>NULL</b> |
|             | carAvailability | 1            | idx_avail_from | 1                   | availability_from | A         | 8393        | <b>NULL</b> | <b>NULL</b> |
|             | carAvailability | 1            | idx_avail_till | 1                   | availability_till | A         | 8394        | <b>NULL</b> | <b>NULL</b> |

query\_block #1

Justification: For this query, when we ran the EXPLAIN ANALYZE command to see execution times, we did not see any significant differences between the times. Sometimes there was only a difference of 0.01 seconds between the different index designs that we chose. Moving forward, we decided to use the visual tool in order to compare query costs. We found that by using the indexes “car\_id,

availability\_id, availability\_till” we got the lowest query cost. Therefore, we will be using this indexing design in our implementation for the next stages.

### Query 1:

- Default:

	Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	N
►	carTripHistory	0	PRIMARY	1	journey_id	A	4123	<a href="#">NULL</a>	<a href="#">NULL</a>	
	carTripHistory	1	fk_lesser_idx	1	lesser_id	A	742	<a href="#">NULL</a>	<a href="#">NULL</a>	
	carTripHistory	1	fk_car_idx	1	car_id	A	135	<a href="#">NULL</a>	<a href="#">NULL</a>	

Query cost: 1863.10

query\_block #1

- Index Design #1: card\_id, owner\_id, VIN

	Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed
►	car	0	PRIMARY	1	car_id	A	213	<a href="#">NULL</a>	<a href="#">NULL</a>
	car	1	fk_owner_idx	1	owner_id	A	168	<a href="#">NULL</a>	<a href="#">NULL</a>
	car	1	idx_vin	1	VIN	A	213	<a href="#">NULL</a>	<a href="#">NULL</a>

Query cost: 1863.10

query\_block #1

- Index Design #2:

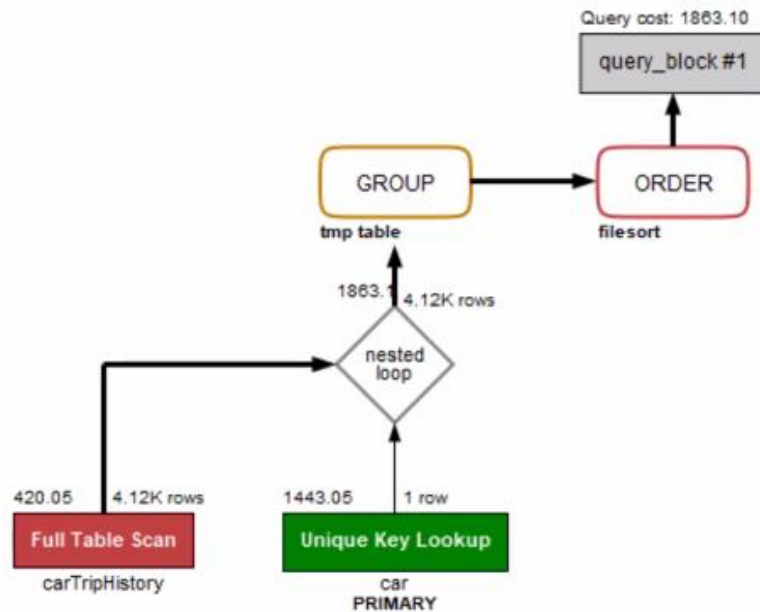
	Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed
►	carTripHistory	0	PRIMARY	1	journey_id	A	4123	<a href="#">NULL</a>	<a href="#">NULL</a>
	carTripHistory	1	fk_lesser_idx	1	lesser_id	A	742	<a href="#">NULL</a>	<a href="#">NULL</a>
	carTripHistory	1	fk_car_idx	1	car_id	A	135	<a href="#">NULL</a>	<a href="#">NULL</a>
	carTripHistory	1	idx_pickup	1	pickup_datetime	A	4002	<a href="#">NULL</a>	<a href="#">NULL</a>

Query cost: 1863.10

query\_block #1

- Index Design #3:

Result Grid								
	Filter Rows:		Export:		Wrap Cell Content:			
Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed
carTripHistory	0	PRIMARY	1	journey_id	A	4123	HULL	HULL
carTripHistory	1	fk_lesser_idx	1	lesser_id	A	742	HULL	HULL
carTripHistory	1	fk_car_idx	1	car_id	A	135	HULL	HULL
carTripHistory	1	idx_dropoff	1	dropoff_datetime	A	4048	HULL	HULL



Justification: For this query, when we ran the EXPLAIN ANALYZE command to see execution times, we did not see any significant differences between the times, nor did we see any differences between the query costs. This behavior persisted despite trying numerous different indexing designs. We provided a visual explain result as context. One possible explanation for this behavior is that our query already runs in the most optimal manner due to our primary key choice (journey\_id).